

# **Data Science & Analysis Foundations with R and Python — Deep Understanding Guide**

Stage 1: Foundations — Statistics, R & Python Basics, Visualisation, Exercises

Author: Maingo Israel

## **Table of Contents**

1. Introduction to Data Science
  2. Core Statistics Refresher
  3. R Fundamentals
  4. Python Fundamentals
  5. Visualization (R & Python)
  6. Practice Exercises
  7. Mini Projects & Suggested Datasets
- Appendix: Useful Resources and Next Steps

# 1. Introduction to Data Science

Data science is the process of extracting knowledge and insights from data. It combines elements of statistics, programming, domain expertise, and communication. The typical workflow is: define the question, gather data, clean and preprocess, explore (EDA), model, evaluate, and communicate results.

In this guide we cover foundational statistics and programming skills in both R and Python, provide code examples, and give exercises that use finance, health, and general datasets. Aim for deep understanding: for each method, ask "why it works", "what assumptions it makes", and "how to check those assumptions".

## 2. Core Statistics Refresher

Descriptive statistics summarize data: mean (average), median (middle), variance and standard deviation (spread), skewness (asymmetry), and kurtosis (tailedness). Use visual and numeric summaries together.

Probability: random variables, expectation, variance. Common distributions: Normal, Student-t, Poisson, Binomial. The

Central Limit Theorem explains why sample means tend toward Normal when sample size is large.

Estimation and hypothesis testing: point and interval estimates, confidence intervals, p-values, and statistical power.

Understand Type I and Type II errors. Always interpret tests within context and effect sizes.

Resampling methods: bootstrap (for estimating sampling distributions) and cross-validation (for model evaluation). These methods are powerful when analytic solutions are hard or model assumptions are uncertain.

Code example:

```
# R: quick descriptive stats
x <- c(1.2, 3.4, 2.5, 4.0, 5.1)
mean(x); median(x); sd(x); var(x)
# Python equivalent
import numpy as np
x = np.array([1.2, 3.4, 2.5, 4.0, 5.1])
x.mean(), np.median(x), x.std(ddof=1), x.var(ddof=1)
```

### 3. R Fundamentals

R is a language for statistics and data analysis. Start with RStudio as an IDE. Key data types: numeric, integer, character, logical, factor, Date. Important data structures: vector, matrix, data.frame, list, tibble.

The tidyverse (dplyr, tidyr, ggplot2) provides consistent data manipulation and plotting grammar. Learn piping (%>%) to chain operations for readable code.

Best practices: comment code, use meaningful variable names, save scripts and RMarkdown for reproducible reports, and use version control (Git).

Code example:

```
# Install tidyverse
install.packages("tidyverse")
library(tidyverse)

# Read CSV and basic EDA
df <- read.csv("data.csv")
glimpse(df)
summary(df)
df %>% select(Date, Close) %>% mutate(Return = diff(log(Close))) %>% head()

# Simple ggplot
library(ggplot2)
ggplot(df, aes(x=Date, y=Close)) + geom_line() + labs(title="Price over time")
```

## 4. Python Fundamentals

Python is a general-purpose language widely used in data science. Use Anaconda and Jupyter or VS Code. Core libraries:

pandas (dataframes), numpy (numerical operations), matplotlib/seaborn/plotly (visualization), scikit-learn (machine learning), statsmodels (statistical models).

Pandas DataFrame is central: `read_csv`, `head()`, `info()`, `describe()`, `groupby()`, `merge()`, `pivot_table()`.

Practice indexing and boolean masks—they're essential.

Best practices: virtual environments, notebooks for exploration, scripts/modules for production, docstrings, and tests for critical code.

Code example:

```
# Python: basic pandas usage
import pandas as pd
import numpy as np

df = pd.read_csv("data.csv", parse_dates=['Date'])
df.info()
df['Return'] = np.log(df['Close']).diff()
df[['Date', 'Close', 'Return']].head()

# Simple matplotlib plot
import matplotlib.pyplot as plt
plt.figure()
plt.plot(df['Date'], df['Close'])
plt.title("Price over time")
plt.xlabel("Date")
plt.ylabel("Close")
plt.show()
```

## 5. Practical Visualization (R & Python)

Visualization is essential for EDA. Use plots to reveal trends, seasonality, outliers, and relationships. Line plots are for time series; histograms/density plots for distribution; scatter plots for relationships; boxplots for spread and outliers.

When plotting time series, always check frequency, missing dates, and annotate events. Combine plots with summary statistics.

Code example:

```
# R: time series plot with ggplot2
library(ggplot2)
df$Date <- as.Date(df$Date)
ggplot(df, aes(x=Date, y=Return)) + geom_line() + geom_smooth(method='loess') + labs(title="Log returns")

# Python: histogram + KDE
df['Return'].hist(bins=50)
df['Return'].plot(kind='kde')
```

## 6. Practice Exercises

Exercises are designed to force active learning. For each exercise, write code in both R and Python and explain results in one paragraph.

Exercise 1: Compute descriptive statistics (mean, median, std, skewness, kurtosis) for daily returns of a chosen stock or commodity. Plot histogram and time series.

Exercise 2: Use bootstrap to create a 95% CI for the mean return. Explain the bootstrap procedure and interpret the interval.

Exercise 3: Load a public health dataset (e.g., patient wait times). Clean missing values, compute summaries by group, and visualize differences.

Exercise 4: Split a dataset into train/test, fit a simple linear regression and report RMSE. Then compare with a tree-based model (random forest).



## 7. Mini Projects & Suggested Datasets

Mini Project A (Finance): Exploratory Analysis of JSE ALSI (or alternative index). Compute log returns, plot volatility, fit ARIMA and simple GARCH(1,1) in R (rugarch) or Python (arch package). Compare models and write concise interpretation.

Mini Project B (Health): Analyze hospital admission times. Identify peak hours/days, model counts with Poisson or negative binomial regression, and visualize results.

Mini Project C (Economic Indicators): Compare GDP growth and unemployment across countries; create small dashboards with plotly or R Shiny.

Suggested datasets: Kaggle, Yahoo Finance (CSV), World Bank (CSV), UCI repository, South African data portals (data.gov.za).

## Appendix: Useful Resources and Next Steps

Appendix: Recommended textbooks and online courses: 'An Introduction to Statistical Learning' (ISLR), 'Elements of Statistical Learning' (Hastie et al.), Andrew Ng's ML course, 'Python for Data Analysis' (Wes McKinney). For time series: 'Time Series Analysis' (Box & Jenkins), 'Forecasting: principles and practice' (Hyndman & Athanasopoulos).

Next steps (after Stage 1): Stage 2 — Data Wrangling & Visualization deep dive; Stage 3 — Machine Learning fundamentals; Stage 4 — Time Series & GARCH; Stage 5 — Projects & Portfolio.

## Appendix: Quick Reference - Selected Code Snippets

R: tidy data and dplyr

```
library(tidyverse)
df %>% drop_na() %>% mutate(Return = log(Close) - lag(log(Close))) %>% filter(!is.na(Return)) %>% summar
```

Python: pandas groupby and pivot

```
# groupby example
df.groupby('Country')['GDP'].mean().sort_values(ascending=False).head()

# pivot example
df.pivot_table(index='Year', columns='Country', values='GDP')
```

R: simple ARIMA with forecast

```
library(forecast)
tsr <- ts(df$Return, frequency=252) # approx trading days
fit <- auto.arima(tsr)
forecast::forecast(fit, h=10)
```

Python: simple scikit-learn regression

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
X = df[['Feature1', 'Feature2']].fillna(0)
y = df['Target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression().fit(X_train, y_train)
mean_squared_error(y_test, model.predict(X_test), squared=False)
```

## Final notes

Study actively: read, code, and explain results in writing. Use both R and Python for the same tasks to build transferable intuition.

Keep a GitHub repo with notebooks, README, and clear explanations. Include small reports (RMarkdown or Jupyter Notebooks) for each mini project.